

Harford the York University MRI Analysis Server

Any sufficiently advanced technology is
indistinguishable from magic.

Hazards of Prophecy: The Failure of Imagination

ARTHUR C. CLARKE

Joseph D. Viviano
Department of Biology
York University
Toronto, ON, CA
joseph.d.viviano@gmail.com

Contents

1	Introduction	2
2	Connecting	2
2.1	Samba: Finder access in OS X	2
2.2	SSH	2
2.3	x2go	3
2.3.1	Troubleshooting x2go	4
2.4	Passwords	4
3	Hardware	5
4	Server Architecture	5
4.1	Partitions	5
4.2	Operating System	6
4.3	Directories	6
4.4	Permissions	6
4.4.1	RAW	7
4.4.2	WORKING	7
4.4.3	ANALYSIS	7
4.4.4	Privacy Notes	8
4.5	Security	8
4.6	Backup	8
5	Software & Configuration	8
5.1	MRI Tools	9
5.2	Programming Languages	9
6	The Pipeline	9
7	Further Reading	9

1 Introduction

This server was built by a complete stranger to server-building, and I undoubtedly made some silly mistakes along the way. This document is my effort to crystallize these decisions for those who paid for it (henceforth *the money*) and whoever takes over responsibility from me (henceforth *the unlucky*). Hopefully this isn't my last job ever, and I still walk this earth. If that is true, feel free to contact me at any time joseph.d.viviano@gmail.com

2 Connecting

To facilitate easy access to the server from your personal computer, on or off campus, we have configured three standard access methods.

2.1 Samba: Finder access in OS X

First, click on your desktop. In your menu bar, select `go, connect to server`, and type `smb://130.63.40.171`, after which you need only to enter your username and password for the server. All non-MRI data will now be accessible and editable through the finder window, assuming you have permission to do so. Contact an administrator to change permissions.

Note to administrators: samba is annoying and maintains its own set of passwords and permissions. This is particularly annoying because you will need to find an elegant way for users to set this password up. I haven't.

First, you need to add an existing user to samba using `smbpasswd -a [username]`. Next we add `[username]` to the `/etc/samba/smbusers` file in the form `[username] = "[username]"`. I've configured samba to mount `/srv/CNNLAB` as a writable directory structure in `/etc/samba/smb.conf` under the `[CNNLAB]` section. More directories could be added in this way.

2.2 SSH

If you are comfortable with terminal-only access, you can use the `ssh` command on all Mac or Linux systems (normally, this is already installed). Windows users should install PuTTY to gain access to the `ssh` command in the Windows terminal.

The major advantage of terminal-only access is lag-free operation on very slow network connections, and the ability to log in to the server from any computer with zero configuration. To access the server, enter the command `ssh -p 45100 [your username]@130.63.40.171`, where `[your username]` is the login name registered on Harford. If you require a username, please contact one of the administrators. You will next need to enter your case-sensitive password associated with your user account and should have access to the server.

2.3 x2go

Normally, you will want to view a full graphic-user interface while you interact with the server. This will allow you to easily view the data remotely, and make use of various programs not available with terminal-only access. To do this, we use x2go ([download x2go client here](#)). On Windows and Mac, you should be able to open x2go by clicking on the installed icon. On Linux, you may need to type `x2goclient` into the terminal.

Configuring x2go is fairly straightforward. First, click the ‘new session’ button, and set these values:

- Session name: [whatever you want ... I chose Harford]
- Host: 130.63.40.171
- Login: [your username]
- SSH port: 45100
- Session type: KDE (or XFCE on Windows or Linux only!)

You might notice that this is almost identical to the information that we use to `ssh` into the server, because that is what we are actually doing! The only additional setting regards the `session type`, as we now need to specify the graphic user interface.

A quick note on KDE & XFCE. If you are on a Macintosh computing device, and you try to use XFCE, your keyboard will not work properly. Therefore, you must select KDE.

The rest of the items can remain blank. You do not need to configure anything specifically under the connection tab. However, here you can set performance values to your liking, especially to compensate for a slow connection. Try moving the slider to the left before changing the picture quality settings manually.

Finally, you can customize session settings such as screen resolution (which will greatly affect your speed), sound, and other misc. settings. These values can be safely left at their defaults. Now you can click 'OK' and log into the server by clicking on the newly created button on the right hand side. You will need to enter your password for the server here.

2.3.1 Troubleshooting x2go

In some cases, you will receive an error involving the RSA key, and will not be allowed to log into the server. Your RSA keys are stored in `~/.ssh/known_hosts`. You can inspect the contents of this file by typing `cat ~/.ssh/known_hosts` into the terminal. If you find an entry that starts with `130.63.40.171`, remove it using either: `gedit ~/.ssh/known_hosts` or `vi ~/.ssh/known_hosts` ([here are some instructions for vi](#)). If you are having trouble doing this, you can simply reset the `known_hosts` file by removing it with `rm ~/.ssh/known_hosts` (this might affect other systems relying on RSA keys, but I will assume it won't if you are following these instructions). You should now be able to log into the server using x2go.

2.4 Passwords

Your default password can be changed to something more easily remembered by doing the following:

1. SSH into the server using port `45100`.
 - Harford uses a non-standard SSH port to prevent random attacks. This requires you to specify it at the command line with the following command: `ssh -p 45100 [your username]130.63.40.171`.
2. Use the `passwd` command.
 - If you type `passwd` at the command line, you will be prompted for your old annoying password, and asked to enter a new one. Your new password must satisfy the following conditions:
 - 1 Uppercase letter
 - 1 lowercase letter
 - 1 number
 - 1 symbol (i.e., `*`, `_`, `#`, `$`, `%`)
 - 12 characters length minimum.

- This does not need to be hard to remember! Computer no think like human.
 - #sQ_ is easy for a computer to guess.
 - Br4in-fun-4eva! is much harder.

Be sure to write your password down. If you forget your password, an administrator can reset it for you. If you *are* an administrator, you should know you can change `username`'s password by dropping to `root` and typing `passwd [username]`.

Also, recall that samba and unix passwords are *not* automatically synchronized. [This really should be fixed.]

3 Hardware

Harford is a `Thinkserver TS440 (8×3.5") HDD Hot-Swappable` server loaded with 4 HDD at the moment. Hard drive caddy part # `03X3969, FRU HS 3.5'' HDD Tray V3.0`. It is loaded with 4 × Western Digital RED 4.0 TB 5400 HDD, configured in RAID 1+0. This currently leaves us with 8 TB of usable hard drive space, striped arrays, and on-site redundancy. It can easily be expanded to 2× that with an additional 4 drives.

3.1 Hardware List

Currently, Harford contains:

1. Thinkserver TS440 [70AQ000CUX]
2. 20 GB RAM [2×8 GB added, 2×2 GB stock.]

4 Server Architecture

4.1 Partitions

At the moment, Harford consists of three partitions:

- (a) `/boot`: a 200 MB partition that is integral to the life of the server. This contains the *Linux kernels* of the system – the server's brain-stem. Do not mess with this partition.

- (b) `/`: a 75 GB partition containing all user `/home` folders, software, and the operating system. This *should* be large enough for indefinite expansion, if people don't store data in their home folders.
- (c) `/srv`: a 6.59 TB partition containing all of our data.

NB: this means that the lion's share of the disk space is found under `/srv` and therefore large files should *always* be kept there!

4.2 Operating System

I chose Ubuntu Server 12.04 LTS for its excellent support, modern features, and compatibility with the NeuroDebian project (neuro.debian.net). I named the server Harford after the Kubrickian hero. When I first installed the server, I was left only with a basic terminal. To get things normal-looking, I had to `sudo apt-get install` the following:

- `openssh-server`
- `xfce4`
- `kde-plasma-desktop`
- `synaptic`
- `lightdm-gtk-greeter`
- `jockey-gtk`
- `dmz-cursor-theme`
- `xubuntu-icon-theme`
- `elementary-icon-theme`

This will produce a basic desktop environment for you to work in. I include both the `xfce4` and `kde` desktop environments, which each have their own merits and utility. They essentially control the graphic user interface of the server but do not differ substantially in actual usability.

A note on what these basic programs do. `openssh-server` enables SSH access through the terminal and is how people can control the server remotely. `xfce4` and `kde-plasma-desktop` are two alternative graphic user interfaces for the computer. `synaptic` allows one to probe the Internet for possible software installations and largely automates that process. Most of the software on the server can be installed and/or uninstalled using this simple program. `lightdm-gtk-greeter` presents the user with a login screen. The remaining installs are all graphic user interface niceties.

4.3 Directories

4.4 Permissions

The MRI data is separated into 3 main branches: RAW, WORKING, and ANALYSIS. These different tiers are more or less editable by various users of the server to strike a balance between data-security and usability. As a general rule, we also don't let unauthorized people look at data they aren't supposed to for both privacy and competitive reasons. These permissions are maintained by a `root`-owned nightly CRON job `maintain_permissions.sh`.

The `grandvizier` user should be used by the system administrators to preform various tasks without the need of dropping to `root`, which I am trying to discourage as it can be dangerous to spend too much time with so much power. If you can see a file owned by the `grandvizier`, that typically means it is being protected.

4.4.1 RAW

These are DICOM files. Mostly used for archival purposes. Shouldn't be edited.

```
owner = grandvizier, rwx
group = staff , r-x
else = --
```

4.4.2 WORKING

These are the files manipulated by the pipeline code. Generally, these files should only be accessed by the pipeline and not manually. Right now, experiment specific group-wise permissions allow for you to go in and delete everything *except* the input RAW data at the bottom of the WORKING tree. This allows you to 'reset' problem subjects or whole users using the `cleanup_X.sh` programs, or manually remove problem files.

```
owner = grandvizier, rwx
group = [experiment], rwx (except inputs which are r-x)
else = , --
```


4.4.3 ANALYSIS

These folders are where people *are* allowed to mess around. Generally, outputs from the pipeline can be copied into the ANALYSIS tree for manipulation. This is zero-risk as there is always an identical copy of the pre-processed files in the WORKING directory.

```
owner = grandvizier, rwx
group = [exrperiment], rwx
else = , --
```

4.4.4 Privacy Notes

The `grandvizier` user is special, and shared among *the money* and *the unlucky*. The `grandvizier` is also capable of destroying millions of taxpayer dollars in a one-line command, so it shouldn't be used by anyone unless required.

All pipeline code will be run by the individuals within an experiment group – and the code will work so long as the individual is permitted to interact with a given data set. This is safe because the pipeline code isn't editable by the users in the first place, so we can't do undue harm to our data by mistake.

4.5 Security

4.6 Backup

5 Software & Configuration

The following is a list of the software installed on the server, and if you are lucky, it is even up-to-date.

- [NeuroDebian](#):
 - [FSL 5.0.6](#)
 - [DICOM2NIFTI](#)
 - [AFNI](#)

- FreeSurfer
 - Caret
 - Python: NiBabel
- General Computing:
 - R
 - gedit
 - Terminator
 - LibreOffice
 - Evince
 - Firefox
 - Inkscape
 - Gimp
 - Dia
 - VLC
 - Samba
- Python Packages:
 - NumPy
 - SciPy
 - Pandas
 - matplotlib
 - scikit-learn
 - scikit-image
 - NetworkX
 - PyMVPA
 - iPython
- Misc – Installed Manually in `/opt`:
 - Gephi
 - MATLAB Compiler Runtime
 - McRetro

5.1 MRI Tools

5.2 Programming Languages

6 The Pipeline

For the pipeline to work, your RAW data must be converted to the appropriate format in the WORKING directory.

7 Further Reading